

MAV - Medical Access & Vision

Intelligent Clinical Decision Support System



Technical Documentation

Version 1.0.0

Project Type:	Clinical Decision Support System
Platform:	Web Application
Backend:	Python Flask
Frontend:	React 18 (Vite)
Database:	MongoDB
AI Engine:	Google Gemini

Authors

Ostello Angelo, Iacovone Vincenzo, Di Maria Matteo

January 2026

Contents

Abstract	4
1 Introduction	5
1.1 Overview and Vision	5
1.1.1 Problem Statement	5
1.1.2 Solution Approach	6
1.2 Key Features	6
2 Technical Architecture	7
2.1 System Overview	7
2.1.1 The Gateway Concept	7
2.1.2 Architecture Diagram	7
2.2 Data Flow Pipeline	8
2.3 Component Interaction	8
2.3.1 Frontend-Backend Communication	8
2.3.2 CORS Configuration	9
3 Design Patterns Analysis	10
3.1 Implemented Patterns	10
3.1.1 Strategy Pattern (Partially Implemented)	10
3.1.2 Decorator Pattern (Implemented)	10
3.2 Conceptual Patterns	11
3.2.1 Chain of Responsibility	11
3.2.2 Observer Pattern	11
3.2.3 Facade Pattern	11
3.3 Summary	11
4 Technology Stack	12
4.1 Backend Technologies	12
4.1.1 Python Flask Framework	12
4.1.2 AI Integration	12
4.1.3 Data Processing Libraries	12
4.1.4 Database Connectivity	13
4.2 Frontend Technologies	13
4.2.1 React 18 with Vite	13
4.2.2 State Management	13
4.2.3 Styling	13
4.3 Infrastructure	13

4.3.1	Database	13
4.3.2	Containerization	14
4.3.3	Dependencies Summary	14
5	Core Components	15
5.1	Patient Management	15
5.2	Doctor Authentication	15
5.2.1	Doctor ID Generation Algorithm	15
5.2.2	Admin-Controlled Registration	16
5.2.3	Authentication Features	16
5.3	Clinical Records	17
5.3.1	Vital Signs Reference	17
5.4	PDF Export	17
6	AI and Diagnostics Engine	18
6.1	Overview	18
6.2	Multimodal Capabilities	18
6.3	Structured Output	18
6.4	Error Handling	19
6.5	Medical Chatbot	19
7	Security and Privacy	20
7.1	Authentication System	20
7.1.1	Access Control and Registration Security	20
7.1.2	User Management System	20
7.1.3	Doctor Registration	22
7.1.4	Session Management	22
7.2	Data Protection	22
7.2.1	Data at Rest	22
7.2.2	Data in Transit	22
7.2.3	Anonymization	22
7.3	Audit Logging	22
7.4	Compliance	23
8	Database Architecture	24
8.1	MongoDB Collections	24
8.2	Indexing Strategy	24
8.3	Connection Configuration	24
9	Deployment Guide	25
9.1	Docker Deployment	25
9.1.1	Quick Start	25
9.2	Manual Deployment	25
9.2.1	Backend	25
9.2.2	Frontend	25
9.3	Environment Variables	25
9.4	Troubleshooting	25

10 Future Improvements	27
10.1 Enhanced Security	27
10.2 Native Applications	27
10.3 OCR Integration	27
10.4 Additional Features	28
10.5 Implementation Roadmap	28
11 Conclusion	29
11.1 Summary	29
11.2 Technical Excellence	29
11.3 Future Vision	29
A API Reference	31
A.1 Principali Endpoint	31

Abstract

MAV is an advanced, production-ready software solution designed to bridge the gap between raw clinical data and modern Artificial Intelligence. This comprehensive system serves as an intelligent routing and processing engine that ingests multi-modal data (text, vital signs, and medical images), validates the information, enriches it with calculated metadata, and leverages state-of-the-art AI to assist healthcare professionals in early disease diagnosis.

The platform addresses critical challenges in modern healthcare environments: data fragmentation through unified data source management, latency through real-time decision support, privacy through strict HIPAA/GDPR compliance mechanisms, and reliability through graceful failure handling with fallback mechanisms.

This documentation provides a complete technical overview of the system architecture, implementation details, deployment procedures, and future improvement opportunities for the MAV platform.

Chapter 1

Introduction

1.1 Overview and Vision

The healthcare industry faces unprecedented challenges in managing the ever-increasing volume of clinical data while maintaining high standards of patient care. MAV emerges as a sophisticated solution to address these multifaceted challenges, providing healthcare professionals with an intelligent platform that not only manages patient records but also offers AI-powered diagnostic assistance.

1.1.1 Problem Statement

Modern healthcare systems struggle with several critical issues:

1. **Data Fragmentation:** Patient information is often scattered across multiple systems, making it difficult to obtain a comprehensive view of a patient's medical history.
2. **Decision Latency:** Healthcare professionals need rapid access to relevant information and analytical insights, especially in emergency triage situations.
3. **Privacy Concerns:** The integration of cloud-based AI services requires stringent data protection mechanisms to comply with regulations such as HIPAA and GDPR.
4. **System Reliability:** Medical systems must handle failures gracefully without compromising patient care continuity.

1.1.2 Solution Approach

MAV addresses these challenges through a carefully designed architecture that implements:

- A unified gateway pattern for clinical data processing
- Real-time AI-powered diagnostic assistance using Google Gemini
- Robust authentication and authorization mechanisms
- Comprehensive audit logging for accountability
- Flexible deployment options via Docker containerization

1.2 Key Features

The platform offers a comprehensive feature set designed to enhance clinical workflows:

- **Patient Management:** Complete patient lifecycle management with support for both Italian citizens (using Fiscal Code validation) and foreign patients (automatic ID generation).
- **Clinical Record Management:** Comprehensive clinical record creation, editing, and retrieval with support for vital signs, symptoms, notes, and file attachments.
- **AI-Powered Diagnostics:** Integration with Google Gemini for multimodal analysis of clinical data and medical images.
- **Medical Chatbot:** An AI-powered conversational assistant for medical queries.
- **Professional Reporting:** PDF export capabilities for clinical documentation.
- **Doctor Authentication:** Secure registration and login system with unique mnemonic doctor IDs.

Chapter 2

Technical Architecture

2.1 System Overview

MAV follows a **Layered Architecture** with a central Gateway component orchestrating the entire data flow pipeline. This architectural approach ensures separation of concerns, maintainability, and scalability.

2.1.1 The Gateway Concept

The core of the backend is the **ClinicalGateway**, which differs fundamentally from a standard CRUD controller. The Gateway treats every incoming clinical request as a payload that must pass through a strict pipeline of handlers, ensuring that no data is ever processed without proper validation, anonymization, and auditing.

2.1.2 Architecture Diagram




Figure 2.1: MAV System Architecture

2.2 Data Flow Pipeline

The clinical data processing follows a well-defined pipeline:

Step 1: Ingestion: The React frontend sends a Patient Record containing symptoms, vital signs, and attachments via HTTP POST request.

Step 2: Gateway Entry: The request enters the ClinicalGateway through the Flask API layer, where initial request parsing occurs.

Step 3: Processing Chain: The request passes through a chain of specialized handlers:

- **ValidationHandler:** Verifies data integrity, including valid Fiscal Codes and realistic vital sign ranges.
- **EnrichmentHandler:** Adds calculated metadata such as patient age from date of birth and BMI calculations.
- **PrivacyHandler:** Anonymizes sensitive fields before external AI processing.
- **TriageHandler:** Calculates initial urgency scores based on configured clinical rules.

Step 4: Strategy Execution: The system selects the appropriate AI strategy (e.g., GeminiStrategy) to analyze the clinical data.

Step 5: Observer Notification: Auditing and Metrics systems record the transaction results for compliance and analytics.

Step 6: Persistence: Validated and enriched data is stored in MongoDB collections.

Step 7: Response: The complete, analyzed result is returned to the Frontend for display.

2.3 Component Interaction

2.3.1 Frontend-Backend Communication

The frontend communicates with the backend through a RESTful API interface. All requests include credentials for session management:

```
1 const response = await fetch(getApiUrl('/api/patient/search'), {  
2   method: 'POST',  
3   headers: { 'Content-Type': 'application/json' },  
4   credentials: 'include',  
5   body: JSON.stringify({ fiscal_code: searchCode })  
6 });
```

2.3.2 CORS Configuration

Cross-Origin Resource Sharing is configured to allow the frontend application to communicate with the backend API:

```
1 CORS(app,
2     resources={
3         r"/api/*": {
4             "origins": allowed_origins,
5             "methods": ["GET", "POST", "PUT", "DELETE", "OPTIONS"],
6             "allow_headers": ["Content-Type", "Authorization"],
7             "supports_credentials": True
8         }
9     })

```

Listing 2.1: CORS Configuration

Chapter 3

Design Patterns Analysis

The system implements several GoF (Gang of Four) design patterns to ensure maintainability, extensibility, and production-ready monitoring capabilities.

3.1 Implemented Patterns

3.1.1 Chain of Responsibility Pattern

The Chain of Responsibility pattern is implemented for the data processing pipeline, handling validation, enrichment, privacy compliance, and triage assessment in a sequential chain.

Architecture

The pipeline consists of four handlers:

- **ValidationHandler:** Validates incoming patient data for required fields
- **EnrichmentHandler:** Enriches data with calculated fields (e.g., age from birth date)
- **PrivacyHandler:** Applies privacy rules and data protection policies
- **TriageHandler:** Performs triage assessment based on clinical data

```
1 class DataProcessingPipeline:  
2     def __init__(self):  
3         self.validation = ValidationHandler()  
4         self.enrichment = EnrichmentHandler()  
5         self.privacy = PrivacyHandler()  
6         self.triage = TriageHandler()  
7  
8         self.validation.set_next(self.enrichment) \  
9             .set_next(self.privacy) \  
10            .set_next(self.triage)  
11  
12     def process(self, data):
```

```
13     return self.validation.handle(data)
```

Listing 3.1: Chain of Responsibility Implementation

Usage

The pipeline is invoked during patient creation to ensure data quality and compliance:

```
1 if data_pipeline:
2     patient_dict = {
3         'fiscal_code': patient.codice_fiscale,
4         'birth_date': patient.data_nascita.isoformat(),
5         'nome': patient.nome,
6         'cognome': patient.cognome
7     }
8     processed_data = data_pipeline.process(patient_dict)
```

Listing 3.2: Pipeline Usage in Patient Creation

3.1.2 Strategy Pattern

The Strategy pattern enables runtime selection of different AI models for medical diagnosis generation. Currently supports Google Gemini with architecture ready for OpenAI GPT and Anthropic Claude.

Available Strategies

- **GeminiStrategy**: Google Gemini AI (active)
- **OpenAIStrategy**: OpenAI GPT (placeholder for future implementation)
- **ClaudeStrategy**: Anthropic Claude (placeholder for future implementation)

```
1 class AIModelContext:
2     def __init__(self, strategy):
3         self._strategy = strategy
4
5     def set_strategy(self, strategy):
6         self._strategy = strategy
7
8     def generate_diagnosis(self, patient_data):
9         return self._strategy.generate_diagnosis(patient_data)
10
11    def get_current_model(self):
12        return self._strategy.get_model_name()
```

Listing 3.3: Strategy Pattern Implementation

Runtime Model Selection

The system initializes with Gemini as the default strategy and allows switching models at runtime:

```

1 if ai_strategy_context:
2     diagnosis_result = ai_strategy_context.generate_diagnosis(
3         patient_data)
4     print(f"Using AI model: {ai_strategy_context.get_current_model()}")
5 else:
6     diagnosis_result = ai_diagnostics.generate_diagnosis(patient_data)

```

Listing 3.4: Strategy Usage in Diagnosis Generation

3.1.3 Observer Pattern

The Observer pattern provides production monitoring through event logging, metrics tracking, and alerting capabilities.

Observer Types

- **AuditLogObserver**: Logs all system events to audit trail database
- **MetricsObserver**: Tracks system metrics (event counts, types, timestamps)
- **AlertObserver**: Sends alerts when error thresholds are exceeded

```

1 class MonitoringSystem(Subject):
2     def __init__(self, db=None):
3         super().__init__()
4         self.audit_observer = AuditLogObserver(db.audit_logs if db else
None)
5         self.metrics_observer = MetricsObserver()
6         self.alert_observer = AlertObserver()
7
8         self.attach(self.audit_observer)
9         self.attach(self.metrics_observer)
10        self.attach(self.alert_observer)
11
12    def log_event(self, event_type, description, user_id=None,
13                  severity='info', metadata=None):
14        event = {
15            'event_type': event_type,
16            'description': description,
17            'user_id': user_id,
18            'severity': severity,
19            'metadata': metadata or {},
20            'timestamp': datetime.now()
21        }
22        self.notify(event)

```

Listing 3.5: Observer Pattern Implementation

Event Monitoring

Events are logged throughout the application lifecycle:

```

1 if monitoring_system:
2     monitoring_system.log_event(
3         'patient_created',
4         f'New patient created: {patient.nome} {patient.cognome}',
5         user_id=session.get('doctor_id'),
6         severity='info',
7         metadata={'codice_fiscale': patient.codice_fiscale}
8     )

```

Listing 3.6: Observer Usage for Event Logging

Metrics Retrieval

Real-time metrics are accessible via API endpoint:

```

1 @app.route('/api/metrics', methods=['GET'])
2 @require_login
3 def get_metrics():
4     if monitoring_system:
5         metrics = monitoring_system.get_metrics()
6         return jsonify({
7             'success': True,
8             'metrics': metrics,
9             'patterns_active': {
10                 'chain_of_responsibility': data_pipeline is not None,
11                 'strategy': ai_strategy_context is not None,
12                 'observer': monitoring_system is not None,
13                 'facade': clinical_facade is not None
14             }
15         })

```

Listing 3.7: Metrics API Endpoint

3.1.4 Facade Pattern

The Facade pattern provides a simplified interface for integration with external clinical systems including HL7/FHIR messaging, PACS, and LIS.

Integrated Systems

- **HL7Interface:** Health Level 7 messaging (ADT, ORM messages)
- **FHIRInterface:** FHIR REST API (Patient, Observation resources)
- **PACSInterface:** Picture Archiving and Communication System
- **LISInterface:** Laboratory Information System

```

1 class ClinicalSystemsFacade:
2     def __init__(self):
3         self.hl7 = HL7Interface()
4         self.fhir = FHIRInterface()
5         self.pacs = PACSInterface()
6         self.lis = LISInterface()
7
8     def register_patient(self, patient_data):
9         results = {
10             'hl7_adt': self.hl7.send_adt_message(patient_data),
11             'fhir_patient': self.fhir.create_patient_resource(
12                 patient_data),
13             'timestamp': datetime.now().isoformat()
14         }
15
16     def submit_clinical_record(self, record_data):
17         results = {}
18         if 'vital_signs' in record_data:
19             results['fhir_observation'] = self.fhir.
20             create_observation_resource(
21                 record_data['vital_signs'])
22         if 'images' in record_data:
23             for img in record_data['images']:
24                 results['pacs_storage'] = self.pacs.store_image(
25                     img.get('data', b''),
26                     img.get('metadata', {}))
27
28     return results

```

Listing 3.8: Facade Pattern Implementation

Simplified Integration

The facade simplifies complex multi-system operations:

```

1 if clinical_facade:
2     clinical_facade.register_patient({
3         'patient_id': patient.codice_fiscale,
4         'nome': patient.nome,
5         'cognome': patient.cognome
6     })

```

Listing 3.9: Facade Usage in Patient Registration

3.1.5 Decorator Pattern

Used for authentication via the `@require_login` decorator that protects sensitive routes by verifying user session presence.

```

1 def require_login(f):
2     @wraps(f)
3     def decorated_function(*args, **kwargs):
4         if 'doctor_id' not in session:
5             return jsonify({'error': 'Unauthorized'}), 401
6         return f(*args, **kwargs)
7     return decorated_function

```

Listing 3.10: Authentication Decorator

3.2 Pattern Implementation Summary

Pattern	Status	Implementation Details
Chain of Responsibility	Implemented	Data processing pipeline with validation, enrichment, privacy, and triage handlers
Strategy	Implemented	AI model selection with Gemini active, OpenAI and Claude ready
Observer	Implemented	Production monitoring with audit logging, metrics tracking, and alerting
Facade	Implemented	Clinical systems integration for HL7/FHIR, PACS, and LIS
Decorator	Implemented	Route authentication and authorization

Table 3.1: Design Pattern Implementation Status

3.3 Benefits and Impact

The implementation of these design patterns provides:

- **Maintainability:** Clear separation of concerns and modular architecture
- **Extensibility:** Easy addition of new AI models, handlers, and observers
- **Monitoring:** Real-time system metrics and event tracking
- **Integration:** Simplified interface for complex external systems
- **Quality:** Consistent data processing and validation pipeline
- **Security:** Centralized authentication and audit logging

Chapter 4

Technology Stack

MAUtilizes a modern, well-integrated technology stack designed for reliability, scalability, and maintainability.

4.1 Backend Technologies

4.1.1 Python Flask Framework

The backend is built using Flask, a lightweight micro-framework that provides flexibility without imposing unnecessary constraints.

Component	Version/Details
Python	3.8+
Flask	3.0.0+
Flask-CORS	4.0.0+

Table 4.1: Core Backend Framework

4.1.2 AI Integration

- **Google Generative AI SDK:** `google-generativeai >= 0.3.0`
- Model: `gemini-3-flash-preview`
- Supports multimodal input (text + images)

4.1.3 Data Processing Libraries

- **NumPy:** Numerical computing
- **Pandas:** Data manipulation and analysis
- **PIL (Pillow):** Image processing for medical attachments
- **PyPDF2:** PDF parsing and generation

- **ReportLab:** Professional PDF report generation

4.1.4 Database Connectivity

- **PyMongo:** MongoDB driver for Python ($>= 4.6.0$)
- SSL/TLS support for secure connections
- Connection pooling for performance

4.2 Frontend Technologies

4.2.1 React 18 with Vite

The frontend leverages modern React development practices:

Technology	Version/Details
React	18.2.0+
React DOM	18.2.0+
Vite	5.0.8+
@vitejs/plugin-react	4.2.1+

Table 4.2: Frontend Framework Stack

4.2.2 State Management

- React Hooks (useState, useEffect)
- Context API for global state
- SessionStorage for data persistence

4.2.3 Styling

- Modern vanilla CSS with Glassmorphism design language
- Responsive layout optimized for dark/light mode
- CSS custom properties for theming

4.3 Infrastructure

4.3.1 Database

- **MongoDB:** NoSQL database for flexible schema management
- MongoDB Atlas for cloud deployment
- Polymorphic clinical data storage

4.3.2 Containerization

- **Docker**: Container runtime
- **Docker Compose**: Multi-container orchestration
- Nginx for frontend static file serving

4.3.3 Dependencies Summary

```
1 # Core
2 flask >=3.0.0
3 flask-cors >=4.0.0
4 pymongo >=4.6.0
5 google-generativeai >=0.3.0
6
7 # Data Processing
8 numpy >=1.21.0
9 pandas >=1.3.0
10
11 # PDF Generation
12 reportlab >=4.0.0
13 PyPDF2 >=3.0.0
14
15 # Security
16 cryptography >=41.0.0
17
18 # Italian Fiscal Code
19 codicefiscale >=0.9
```

Listing 4.1: Key Backend Dependencies

Chapter 5

Core Components

5.1 Patient Management

The module manages the entire patient lifecycle with support for:

- **Italian Citizens:** Identification via validated Fiscal Code
- **Foreign Citizens:** Automatic unique ID generation (format: XX-YYYY)
- Automatic age and metadata calculation
- Allergy and permanent disease management

5.2 Doctor Authentication

The authentication system implements a secure, admin-controlled registration process to prevent unauthorized access and fraudulent doctor ID generation.

5.2.1 Doctor ID Generation Algorithm

Doctor IDs are automatically generated using a cryptographically secure algorithm that creates unique, memorable identifiers:

- **Format:** 6-character alphanumeric code (e.g., MR7X9Z, AB4K2L)
- **Structure:** First letter of first name + First letter of last name + 4 random characters
- **Character Set:** Uppercase letters (A-Z) and digits (0-9)
- **Randomness:** Uses Python's `secrets` module for cryptographically strong random generation
- **Uniqueness:** System checks for duplicates before assignment

Example: A doctor named "Mario Rossi" would receive an ID like "MR7X9Z" where:

- M = First letter of "Mario"
- R = First letter of "Rossi"
- 7X9Z = Randomly generated secure suffix

5.2.2 Admin-Controlled Registration

To prevent unauthorized access and fraudulent doctor ID creation, the registration functionality is restricted:

- **Production Mode:** Only authenticated administrators can create new doctor accounts
- **Regular Users:** See only the login form requiring existing doctor ID and password
- **Admin Users:** Access both login and registration forms to create accounts for legitimate medical professionals
- **ID Distribution:** Administrators securely share generated IDs with authorized doctors through verified channels

This two-tier access control ensures that:

1. Non-medical personnel cannot obtain fake doctor IDs
2. All doctor accounts are verified and authorized by administrators
3. The system maintains a trusted registry of healthcare professionals
4. Accountability is maintained through controlled account creation

5.2.3 Authentication Features

- SHA-256 password hashing for secure credential storage
- Session management with secure, HTTP-only cookies
- 7-day session lifetime with automatic expiration
- Support for doctor specialization and affiliated hospital tracking
- Login attempt monitoring and rate limiting

5.3 Clinical Records

Clinical records include:

- Visit information (ID, timestamp, priority)
- Chief complaint and symptoms
- Vital signs (blood pressure, heart rate, temperature, saturation)
- Attachments (medical images, documents)
- Clinical notes

5.3.1 Vital Signs Reference

Parameter	Unit	Normal Range
Blood Pressure	mmHg	90/60 - 120/80
Heart Rate	bpm	60 - 100
Temperature	°C	36.1 - 37.2
O2 Saturation	%	95 - 100

Table 5.1: Vital Signs Ranges

5.4 PDF Export

Professional report generation with ReportLab including patient data, clinical records, attached images and AI diagnoses.

Chapter 6

AI and Diagnostics Engine

6.1 Overview

The AI module provides structured diagnostic support to healthcare professionals through integration with Google Gemini.

6.2 Multimodal Capabilities

The system analyzes simultaneously:

- **Structured Clinical Data:** Patient demographics, medical history, current symptoms, vital signs
- **Medical Images:** X-rays, ECGs, dermatological photographs (Base64 format, automatic RGB conversion)

6.3 Structured Output

The AI generates complete clinical assessments including:

1. Clinical data analysis and image interpretation
2. Presumptive diagnosis with probability and clinical reasoning
3. Differential diagnoses with supporting evidence
4. Recommended diagnostic tests
5. Treatment plan (pharmacological and non-pharmacological)
6. Monitoring and follow-up plan
7. Urgency assessment and intervention timeline

6.4 Error Handling

Implementation of retry logic with exponential backoff (max 3 attempts) to ensure service reliability.

6.5 Medical Chatbot

Dedicated AI instance for medical queries with:

- Isolated API key
- Conversational interface
- Strict medical-only policy
- Real-time response generation

Chapter 7

Security and Privacy

7.1 Authentication System

7.1.1 Access Control and Registration Security

The system implements a critical security measure to prevent unauthorized account creation:

- **Admin-Only Registration:** Doctor account creation is restricted to authenticated administrators
- **Public Interface:** Regular users only see the login form, preventing self-registration
- **Verified Accounts:** All doctor accounts must be created by trusted administrators
- **ID Security:** The cryptographically generated doctor IDs cannot be guessed or forged

This approach prevents:

- Unauthorized individuals from creating fake doctor accounts
- Self-registration by non-medical personnel
- Identity fraud and impersonation
- Uncontrolled access to sensitive patient data

7.1.2 User Management System

The platform will include a dedicated administrative management interface for user and role administration:

Management Dashboard Features

- **User Data Management:** View, edit, and manage all registered doctor accounts
- **Role Assignment:** Dynamically assign or modify user roles (Doctor, Admin)
- **Testing Environment:** Secure environment for application testing and validation
- **Account Status Control:** Activate, deactivate, or delete user accounts

Role Management Capabilities

The management system enables administrators to:

1. **Promote to Admin:** Elevate trusted doctors to administrator status for account management duties
2. **Revoke Admin Rights:** Demote administrators back to regular doctor status when necessary
3. **Deactivate Accounts:** Suspend access for retired doctors or those on leave
4. **Remove Accounts:** Permanently delete accounts for struck-off doctors or terminated personnel
5. **Audit Trail:** Track all role changes and account modifications for compliance

Use Cases

Common administrative scenarios handled by the management system:

- **Retirement:** Deactivating accounts for doctors who have retired from practice
- **Disciplinary Actions:** Removing access for doctors who have been struck off the medical register
- **Temporary Leave:** Suspending accounts during sabbaticals or extended leave
- **Admin Rotation:** Managing administrative privileges among trusted staff
- **Emergency Access:** Quickly granting or revoking access in critical situations

This centralized management approach ensures:

- Efficient user lifecycle management
- Rapid response to personnel changes

- Maintained system security and integrity
- Compliance with healthcare regulatory requirements
- Clear accountability through comprehensive audit logging

7.1.3 Doctor Registration

- Unique mnemonic ID generation using `secrets` module
- Password validation (minimum 6 characters)
- SHA-256 hashing
- Duplicate ID prevention

7.1.4 Session Management

Secure cookies with HTTPS-only configuration in production, HttpOnly, SameSite policy and 7-day duration.

7.2 Data Protection

7.2.1 Data at Rest

- PII separation from clinical data in MongoDB
- Encryption options via MongoDB Atlas
- Secure database credentials management

7.2.2 Data in Transit

- Forced HTTPS in production
- TLS/SSL for database connections
- Secure cookie transmission

7.2.3 Anonymization

Anonymization functionality for external processing: removal of identifying data while maintaining relevant clinical information (allergies, diseases).

7.3 Audit Logging

Logging of all significant actions:

- Patient record creation/modification
- Clinical record access
- Export operations
- Authentication events
- AI diagnostic requests

7.4 Compliance

The system is designed considering:

- **HIPAA:** Health Insurance Portability and Accountability Act
- **GDPR:** General Data Protection Regulation
- Audit trail maintenance
- Data minimization

Chapter 8

Database Architecture

8.1 MongoDB Collections

The database uses 5 main collections:

- **doctors**: Doctor credentials and profiles (unique doctor_id)
- **patients**: Patient demographics (unique codice_fiscale/patient_id)
- **patient_records**: Clinical records (unique encounter_id, linked to patient_id)
- **audit_logs**: Immutable system action history
- **chatbot_sessions**: Medical chatbot sessions

8.2 Indexing Strategy

Collection	Field	Type
patients	codice_fiscale	Unique
patients	patient_id	Unique
patient_records	encounter_id	Unique
patient_records	patient_id	Regular
doctors	doctor_id	Unique
audit_logs	timestamp	Regular

Table 8.1: Database Indexes

8.3 Connection Configuration

MongoDB connection with TLS/SSL, configurable timeouts and connection pooling for optimal performance.

Chapter 9

Deployment Guide

9.1 Docker Deployment

9.1.1 Quick Start

```
1 git clone repository
2 cp .env.example .env
3 # Configure environment variables
4 docker-compose up -d --build
```

9.2 Manual Deployment

9.2.1 Backend

```
1 cd backend
2 python -m venv venv
3 venv\Scripts\activate # Windows
4 pip install -r requirements.txt
5 python webapp/app.py
```

9.2.2 Frontend

```
1 cd frontend
2 npm install
3 npm run dev
```

9.3 Render Deployment

Render is a modern cloud platform that provides seamless deployment for web applications with automatic SSL, continuous deployment from Git, and managed infrastructure.

9.3.1 Platform Overview

Render offers several advantages for deploying MAV:

- **Automatic HTTPS:** Free SSL certificates with automatic renewal
- **Git Integration:** Automatic deployments triggered by repository commits
- **Zero-Downtime Deploys:** Rolling updates without service interruption
- **Managed Infrastructure:** No server management required
- **Environment Variables:** Secure configuration management through web dashboard
- **Custom Domains:** Support for custom domain mapping
- **Health Checks:** Automatic service monitoring and restart

9.3.2 Service Configuration

MAV can be deployed on Render using two separate services:

Backend Service (Web Service)

- **Type:** Web Service
- **Runtime:** Python 3.8+
- **Build Command:** `pip install -r backend/requirements.txt`
- **Start Command:** `python backend/webapp/app.py`
- **Port:** Auto-detected from Flask application
- **Health Check Path:** `/api/health` (recommended)

Frontend Service (Static Site)

- **Type:** Static Site
- **Build Command:** `cd frontend && npm install && npm run build`
- **Publish Directory:** `frontend/dist`
- **Auto-Deploy:** Enabled on main branch commits

9.3.3 Deployment Steps

1. **Repository Connection:** Link GitHub/GitLab repository to Render
2. **Service Creation:** Create separate services for backend and frontend
3. **Environment Configuration:** Set required environment variables through Render dashboard
4. **Build Configuration:** Define build and start commands for each service
5. **Deploy:** Trigger initial deployment manually or via Git push
6. **Domain Setup:** Configure custom domain or use provided `.onrender.com` subdomain

9.3.4 Environment Configuration on Render

Environment variables must be configured in the Render dashboard for each service:

Service	Required Variables
Backend	GEMINI_API_KEY MONGODB_CONNECTION_STRING FLASK_SECRET_KEY FRONTEND_URL
Frontend	VITE_API_URL

Table 9.1: Render Environment Variables by Service

9.3.5 Production Considerations

- **Instance Type:** Select appropriate instance size based on expected traffic
- **Region Selection:** Choose region closest to primary user base for optimal latency
- **Scaling:** Configure auto-scaling policies for backend service
- **Monitoring:** Enable Render's built-in metrics and logging
- **Backup Strategy:** Ensure MongoDB Atlas backups are configured independently

9.3.6 Continuous Deployment Workflow

Render automatically deploys when changes are pushed to the configured branch:

1. Developer pushes code to main branch
2. Render detects commit via webhook

3. Build process executes automatically
4. Health checks verify successful deployment
5. New version goes live with zero downtime
6. Previous version remains available for instant rollback

9.4 Environment Variables

Variable	Description
GEMINI_API_KEY	Google Gemini API key
MONGODB_CONNECTION_STRING	MongoDB Atlas connection string
FLASK_SECRET_KEY	Session encryption key
FRONTEND_URL	Frontend URL (CORS)
VITE_API_URL	Backend API URL

Table 9.2: Required Environment Variables

9.5 Troubleshooting

- **CORS:** Check protocol and trailing slashes
- **AI:** Images under 4MB, standard formats
- **Database:** IP whitelist MongoDB Atlas
- **Sessions:** Cookie configuration cross-domain

Chapter 10

Future Improvements

10.1 Enhanced Security

- Advanced password hashing (bcrypt/Argon2)
- Multi-Factor Authentication (TOTP, SMS, FIDO2)
- JWT-based authentication with token rotation
- Granular RBAC (Admin, Doctor, Nurse, Receptionist)

10.2 Native Applications

- **Desktop:** Electron, Tauri or Flutter Desktop
- **Mobile:** React Native or Flutter
- **PWA:** Service Workers, offline mode, push notifications

10.3 OCR Integration

OCR integration (Tesseract, Google Cloud Vision, Azure) for text extraction from:

- Medical prescriptions
- Laboratory reports
- Discharge letters
- Insurance documents

10.4 Additional Features

- **HL7 FHIR:** Interoperability with other healthcare systems
- **Analytics Dashboard:** Population statistics, disease trends
- **Telemedicine:** Video consultations, secure messaging
- **Enhanced AI:** Multiple providers, specialty-specific models
- **Appointment Scheduling:** Appointment and resource management
- **Prescription Management:** e-Prescription, drug interaction checking
- **Performance:** Redis caching, CDN, query optimization
- **Accessibility:** WCAG 2.1 AA compliance

10.5 Implementation Roadmap

Phase	Feature	Priority
Q1	Enhanced Security, OCR	High
Q2	PWA, HL7 FHIR	Medium
Q3	Desktop App, Analytics	Medium
Q4	Telemedicine, Mobile	Low

Table 10.1: Implementation Roadmap

Chapter 11

Conclusion

11.1 Summary

MAV represents an advanced clinical decision support system, combining modern web technologies with artificial intelligence to assist healthcare professionals.

Key achievements:

- **Robust Architecture:** Gateway pattern with well-defined design patterns
- **AI Diagnostics:** Google Gemini integration for multimodal analysis
- **Patient Management:** Complete management with support for Italian and foreign citizens
- **Security-First:** Authentication, authorization and audit logging
- **Modern Deployment:** Docker containerization

11.2 Technical Excellence

The project demonstrates adherence to best practices:

- Clean code architecture with separation of concerns
- Design pattern implementation
- Error handling with retry mechanisms
- Scalable database design

11.3 Future Vision

The roadmap positions MAV for continuous evolution:

- Enterprise-grade enhanced security

- Native applications for accessibility
 - OCR for paper documents
 - Healthcare interoperability standards
-

Document generated on January 1, 2026

Appendix A

API Reference

A.1 Principali Endpoint

Method	Endpoint	Description
POST	/api/auth/register	Registrazione medico
POST	/api/auth/login	Login medico
POST	/api/patient/search	Ricerca paziente
POST	/api/patient/create	Creazione paziente
POST	/api/record/add	Aggiunta cartella clinica
POST	/api/diagnosis/generate	Generazione diagnosi AI
GET	/api/export/:fiscal_code	Export PDF

Table A.1: API Endpoints Principali